

Lab 11:
Probabilistic Models of
Geographic Range Evolution
By Will Freyman

1 Before you begin

Today's lab will be completed in R. Please install R: <https://www.r-project.org/>

We'll be using the excellent R package **BioGeoBears** which was developed by Nick Matkze, who was a GSI of this course! His website (and the source of the example script we'll use) is here: <http://phylo.wikidot.com/biogeobears>

2 Introduction to probabilistic biogeography models

Probabilistic modeling of geographic range evolution allows us to use the standard set of statistical model choice procedures (e.g. AIC, BIC, likelihood ratios, etc.) to test different biogeographical scenarios, and to make statistically sound inference of ancestral geographic ranges over a phylogeny. For example, we can test whether the observed biogeographic distribution of a clade is best explained with a model that allows for vicariance and long-distance dispersal versus a model that allows for only vicariance.

There are many parsimony based approaches for studying biogeography that will be covered in lecture. The first likelihood approaches to biogeography were presented in the ground breaking Ree and Smith [2008]. We'll discuss these probabilistic models and their many extensions in lecture (also see Figure 1 for a summary), so here we'll focus on applying them to the the phylogeny of the Hawaiian members of the plant genus *Psychotria* published in Ree and Smith [2008].

3 Installing BioGeoBears

We will use BioGeoBears to implement the 2-parameter Dispersal-Extinction-Cladogenesis (DEC) model from Ree and Smith [2008] as well as the DEC+J model that includes a long distance dispersal (jump or founder event speciation) parameter. See Figure 1 for details on the different models. We'll use the Likelihood Ratio Test (LRT) and Akaike information criterion (AIC) to see which model fits the data better.

3.1 Install and setup BioGeoBears

Start up R and be sure to set your working directory so that you can find the output file we will be generating (look back at previous labs if you don't remember how to set the working directory). Install the BioGeoBears package:

```
install.packages("BioGeoBEARS", dependencies=TRUE, repos="http://cran.rstudio.com")
```

Now lets load the package and some dependencies:

```
library(optimx)
library(FD)
library(snow)
```

	Process	Ranges Before After	Character mapping	DIVA	DEC (GeoSSE, LAGRANGE)	BayArea, BBM (RASP)	Parameter of BioGeoBEARS Supermodel
Anagenetic	Dispersal			✓	✓	✓	d (& x, b)
	Extinction			✓	✓	✓	e (& u, b)
	Range-switching		✓				a (& x, b)
Cladogenetic	Sympatry (narrow)		✓	✓	✓	✓	y (& $mx0/y$)
	Sympatry (widespread)					✓	y (& $mx0/y$)
	Sympatry (subset)				✓		s (& $mx0/s$)
	Vicariance (narrow)			✓	✓		v (& $mx0/v$)
	Vicariance (widespread)			✓			v (& $mx0/v$)
	Founder						j (& $x, mx0/j$)

Figure 1: The processes assumed by different historical biogeography methods. Each of these processes is controlled by the specified parameter(s) in the BioGeoBEARS supermodel, allowing them to be turned on or off, or estimated from the data. Note that whether or not the data support a particular free parameter is an empirical question that should be tested with model choice procedures. Note also that this graphic deals only with the range-changing processes assumed by the different methods. BioGeoBEARS does not attempt to replicate e.g. the parsimony aspect of DIVA, just the processes allowed by DIVA (the BioGeoBEARS “DIVA” model could be called “DIVALIKE” to emphasize that it is a likelihood implementation of the processes assumed by DIVA). Text and image from http://phylo.wikidot.com/biogeobears#BioGeoBEARS_supermodel_graphic.

```
library(parallel)
library(BioGeoBEARS)
```

Now let's load some of the source files necessary to run BioGeoBEARS:

```
source("http://phylo.wdfiles.com/local--files/biogeobears/cladoRcpp.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_add_fossils_randomly_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_basics_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_calc_transition_matrices_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_classes_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_detection_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_DNA_cladogenesis_sim_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_extract_Qmat_COOmat_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_generics_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_models_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_on_multiple_trees_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_plots_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_readwrite_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_simulate_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_SSEsim_makePlots_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_SSEsim_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_stochastic_mapping_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_stratified_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_univ_model_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/calc_uPASS_probs_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/calc_loglike_sp_v01.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/get_stratified_subbranch_top_downpass_likelihooods_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/runBSM_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/stochastic_map_given_inputs.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/summarize_BSM_tables_v1.R")
source("http://phylo.wdfiles.com/local--files/biogeobears/BioGeoBEARS_traits_v1.R") # added traits model
calc_loglike_sp = compiler::cmpfun(calc_loglike_sp_prebyte)
calc_independent_likelihooods_on_each_branch = compiler::cmpfun(calc_independent_likelihooods_on_each_branch_prebyte)
```

And finally, let's get the directory of the example data that installed with BioGeoBEARS:

```
extdata_dir = np(system.file("extdata", package="BioGeoBEARS"))
```

3.2 Load the phylogeny

This is the phylogeny of the Hawaiian members of the plant genus *Psychotria* from Ree and Smith [2008].

```
tree_file_name = np(paste(addslash(extdata_dir), "Psychotria_5.2.newick", sep=""))
tr = read.tree(tree_file_name)
```

Let's plot the tree:

```
plot(tr)
title("Example Psychotria phylogeny from Ree & Smith (2008)")
axisPhylo()
```

3.3 Load the geography data

Now we need to load data on the geographic range of each extant species in our phylogeny.

```
geo_file_name = np(paste(addslash(extdata_dir), "Psychotria_geog.data", sep=""))
tipranges = getranges_from_LagrangePHYLIP(lgdata_fn=geo_file_name)
```

Let's take a look at the geographic range data

```
tipranges
```

Question 1:

Geographic range is being modeled here as a discrete character state, where the character state represents the combination of areas that a lineage inhabits at any given point in time. Models of biogeographic range evolution are essentially no different than the other models of discrete character evolution we have looked at over the course of the semester (with the exception of including cladogenetic events).

1. DNA substitution models have 4 discrete states. How many states will our biogeographic model have? Remember, these models allow for a lineage to be in more than one area at a time.
2. Does it make sense to allow the character state 0 0 0 0? What would this represent?
3. To calculate likelihoods of discrete character evolution models we have to exponentiate the transition matrix, which can be very computationally demanding for large matrices. If we were performing inference on a dataset with 10 areas instead of 4, how many states would our model have? How large would the anagenetic transition rate matrix be?
4. To model cladogenetic change we need an additional transition probability matrix that represents probabilities for all combinations of the state before cladogenesis and after on each of the two daughter lineages. Given 10 areas, how large would the cladogenetic transition probability matrix be?

3.4 Setup the DEC model

First we will run the analysis using the DEC model. The DEC is the default model in BioGeoBEARS, so it is relatively straightforward to setup.

```
BioGeoBEARS_run_object = define_BioGeoBEARS_run()
```

Give BioGeoBEARS the location of the example input files:

```
BioGeoBEARS_run_object$trfn = tree_file_name  
BioGeoBEARS_run_object$geogfn = geo_file_name
```

And let's configure our analysis. If you are running this on your own dataset be sure to adjust the maximum range size (that is the maximum number of areas a lineage can inhabit at an given point – usually this is just the total number of areas in your dataset).

```
BioGeoBEARS_run_object$max_range_size = 4  
BioGeoBEARS_run_object$min_branchlength = 0.000001  
BioGeoBEARS_run_object$include_null_range = TRUE  
BioGeoBEARS_run_object$num_cores_to_use = 1  
BioGeoBEARS_run_object$force_sparse = FALSE  
BioGeoBEARS_run_object = readfiles_BioGeoBEARS_run(BioGeoBEARS_run_object)  
BioGeoBEARS_run_object$return_condlikes_table = TRUE  
BioGeoBEARS_run_object$calc_TTL_loglike_from_condlikes_table = TRUE  
BioGeoBEARS_run_object$calc_ancprobs = TRUE
```

Now we are ready to run the analysis:

```
results_DEC = bears_optim_run(BioGeoBEARS_run_object)
```

Question 2:

Take a look at the `results_DEC` object. What is the maximum likelihood estimate of the rate of anagenetic “dispersal” (range expansion)? And the rate of anagenetic “extinction” (range contraction)?

3.5 Setup the DEC+J model

Now we will run the analysis using the DEC+J model. This include the “jump” parameter for long distance dispersal / founder speciation events.

```
BioGeoBEARS_run_object = define_BioGeoBEARS_run()
```

Give BioGeoBEARS the location of the example input files:

```
BioGeoBEARS_run_object$trfn = tree_file_name
BioGeoBEARS_run_object$geogfn = geo_file_name
```

Just like before, let’s configure the analysis. If you are running this on your own dataset be sure to adjust the maximum range size. These settings are all the same as they were for the DEC model above.

```
BioGeoBEARS_run_object$max_range_size = 4
BioGeoBEARS_run_object$min_branchlength = 0.000001
BioGeoBEARS_run_object$include_null_range = TRUE
BioGeoBEARS_run_object$num_cores_to_use = 1
BioGeoBEARS_run_object$force_sparse = FALSE
BioGeoBEARS_run_object = readfiles_BioGeoBEARS_run(BioGeoBEARS_run_object)
BioGeoBEARS_run_object$return_condlikes_table = TRUE
BioGeoBEARS_run_object$calc_TTL_loglike_from_condlikes_table = TRUE
BioGeoBEARS_run_object$calc_ancprobs = TRUE
```

And now finally, set up DEC+J model. We’ll use the maximum likelihood parameter value estimates from the DEC analysis we already ran to get good starting points for the hill-climbing heuristic.

```
dstart = results_DEC$outputs@params_table["d","est"]
estart = results_DEC$outputs@params_table["e","est"]
```

Set the starting values in our analysis object:

```
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["d","init"] = dstart
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["d","est"] = dstart
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["e","init"] = estart
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["e","est"] = estart
```

The DEC is a 2-parameter model nested within the 3-parameter DEC+J, so we need to add J as a new free parameter to estimate. We also need to assign it an initial value.

```
jstart = 0.0001
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["j","type"] = "free"
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["j","init"] = jstart
BioGeoBEARS_run_object$BioGeoBEARS_model_object@params_table["j","est"] = jstart
```

Now we are ready to run the analysis:

```
results_DECJ = bears_optim_run(BioGeoBEARS_run_object)
```

3.6 Plotting ancestral states

Let's plot the results of the two models and look at the ancestral state reconstructions.

```
pdfn = "Psychotria_DEC_vs_DEC+J.pdf"
pdf(pdfn, width=6, height=6)
```

Plot the DEC ancestral states:

```
analysis_titledtxt = "DEC on Psychotria"
results_object = results_DEC
scriptdir = np(system.file("extdata/a_scripts", package="BioGeoBEARS"))
res2 = plot_BioGeoBEARS_results(results_object, analysis_titledtxt, addl_params=list("j"), plotwhat="text",
                                label.offset=0.45, tipcex=0.7, statecex=0.7, splitcex=0.6, titlecex=0.8, plotsplits=TRUE,
                                cornercoords_loc=scriptdir, include_null_range=TRUE, tr=tr, tipranges=tipranges)
plot_BioGeoBEARS_results(results_object, analysis_titledtxt, addl_params=list("j"), plotwhat="pie",
                          label.offset=0.45, tipcex=0.7, statecex=0.7, splitcex=0.6, titlecex=0.8, plotsplits=TRUE,
                          cornercoords_loc=scriptdir, include_null_range=TRUE, tr=tr, tipranges=tipranges)
```

Plot the DEC+J ancestral states:

```
analysis_titledtxt = "DEC+J on Psychotria"
results_object = results_DECJ
scriptdir = np(system.file("extdata/a_scripts", package="BioGeoBEARS"))
res1 = plot_BioGeoBEARS_results(results_object, analysis_titledtxt, addl_params=list("j"), plotwhat="text",
                                label.offset=0.45, tipcex=0.7, statecex=0.7, splitcex=0.6, titlecex=0.8, plotsplits=TRUE,
                                cornercoords_loc=scriptdir, include_null_range=TRUE, tr=tr, tipranges=tipranges)
plot_BioGeoBEARS_results(results_object, analysis_titledtxt, addl_params=list("j"), plotwhat="pie",
                          label.offset=0.45, tipcex=0.7, statecex=0.7, splitcex=0.6, titlecex=0.8, plotsplits=TRUE,
                          cornercoords_loc=scriptdir, include_null_range=TRUE, tr=tr, tipranges=tipranges)
```

Make a PDF file. We'll see each model's ancestral state reconstructions plotted twice: 1) with the the maximum likelihood estimate of the ancestral ranges, and 2) with pie charts at each node that show the probability of each ancestral state.

```
dev.off()
cmdstr = paste("open ", pdfn, sep="")
system(cmdstr)
```

Question 3:

Compare the estimated ancestral ranges of the lineages leading to *P_hexandra_Oahu* all the way back to the root of the tree. Explain the results in context of biogeographic hypothesis testing. Which hypothesis makes more sense to you given Hawaiian island geography?

3.7 Model testing

Let's use the Likelihood Ratio Test (LRT) and Akaike information criterion (AIC) to see which model fits the data better.

First get the log likelihoods:

```
LnL_2 = get_LnL_from_BioGeoBEARS_results_object(results_DEC)
LnL_1 = get_LnL_from_BioGeoBEARS_results_object(results_DECJ)
```

The AIC is calculated with the log likelihoods and the number of parameters in each model:

```
numparams1 = 3
numparams2 = 2
```

Calculate AIC:

```
stats = AICstats_2models(LnL_1, LnL_2, numparams1, numparams2)
stats$AIC1
stats$AIC2
stats$pval
```

Question 4:

Which model does the AIC support? What about the LRT?

Please email me the following:

1. Your PDF of the ancestral state reconstructions.
2. The answers to questions 1-4.

References

Richard H Ree and Stephen A Smith. Maximum likelihood inference of geographic range evolution by dispersal, local extinction, and cladogenesis. *Systematic Biology*, 57(1):4–14, 2008.